



BROCKMANN  
CONSULT

---

Turning Environmental Data Into Knowledge

# Application packages in the xcube ecosystem

**Pontus Lurcock, Brockmann Consult**

4 September 2025

# Requirements (you should have these already!)

See the course website at  
<https://xcube-dev.github.io/summerschool25/>  
for full details.

- Your preferred **terminal emulator**, preferably running the bash shell.
- A conda-based **Python environment**, which we'll use to set up an environment in which to run the Python tools. I recommend mamba. Installation instructions here.
- git
- docker

# Overview of this course

In this course, you'll get to know:

- OGC Earth Observation **Application Packages**, a framework that lets you package almost any software into a reusable module for deployment on cloud infrastructure.
- **xcube**, a multitalented, ever-growing toolkit and ecosystem for working with data cubes in Python.
- **xcengine**, a new tool that lets you automatically turn your Python Jupyter notebooks into both Application Packages and xcube server containers.

# Course structure

For each part of the course, I'll first present a few slides from this deck to introduce the topic and give background information.

Then we'll move to the hands-on section of that part. In the GitHub repository is a notebook called **summerschool.ipynb** which will guide you through the hands-on parts.

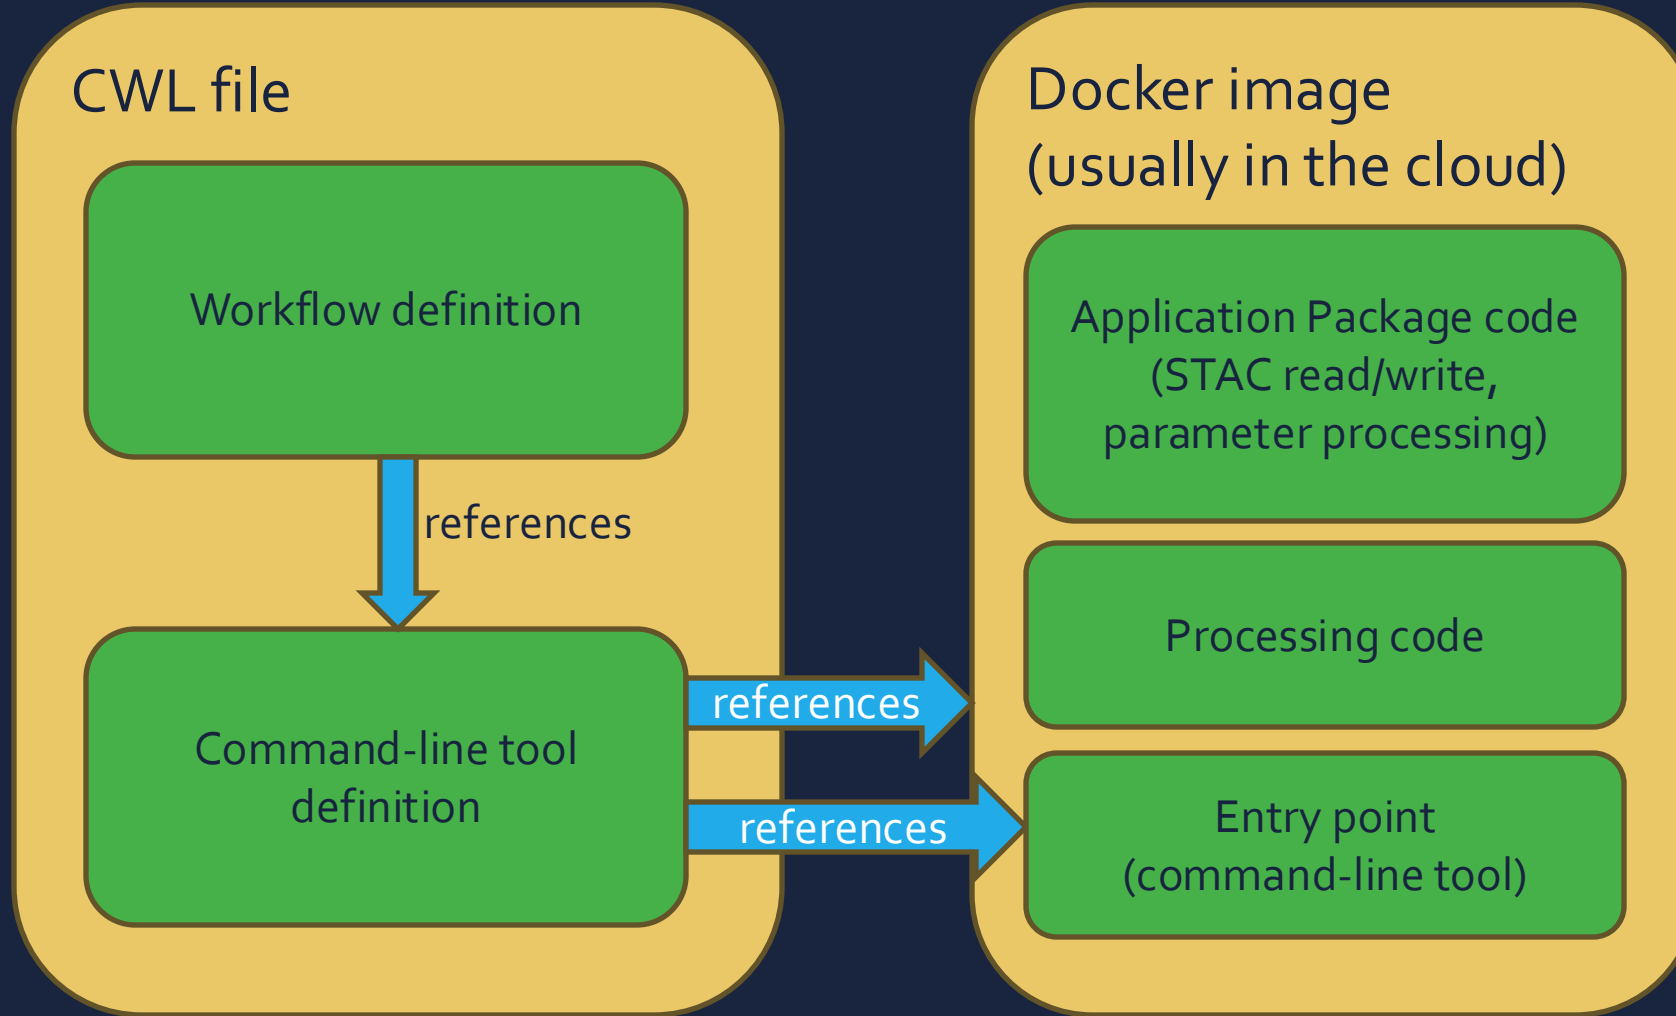


# Part 1: Introducing Application Packages

# OGC EO Application Packages

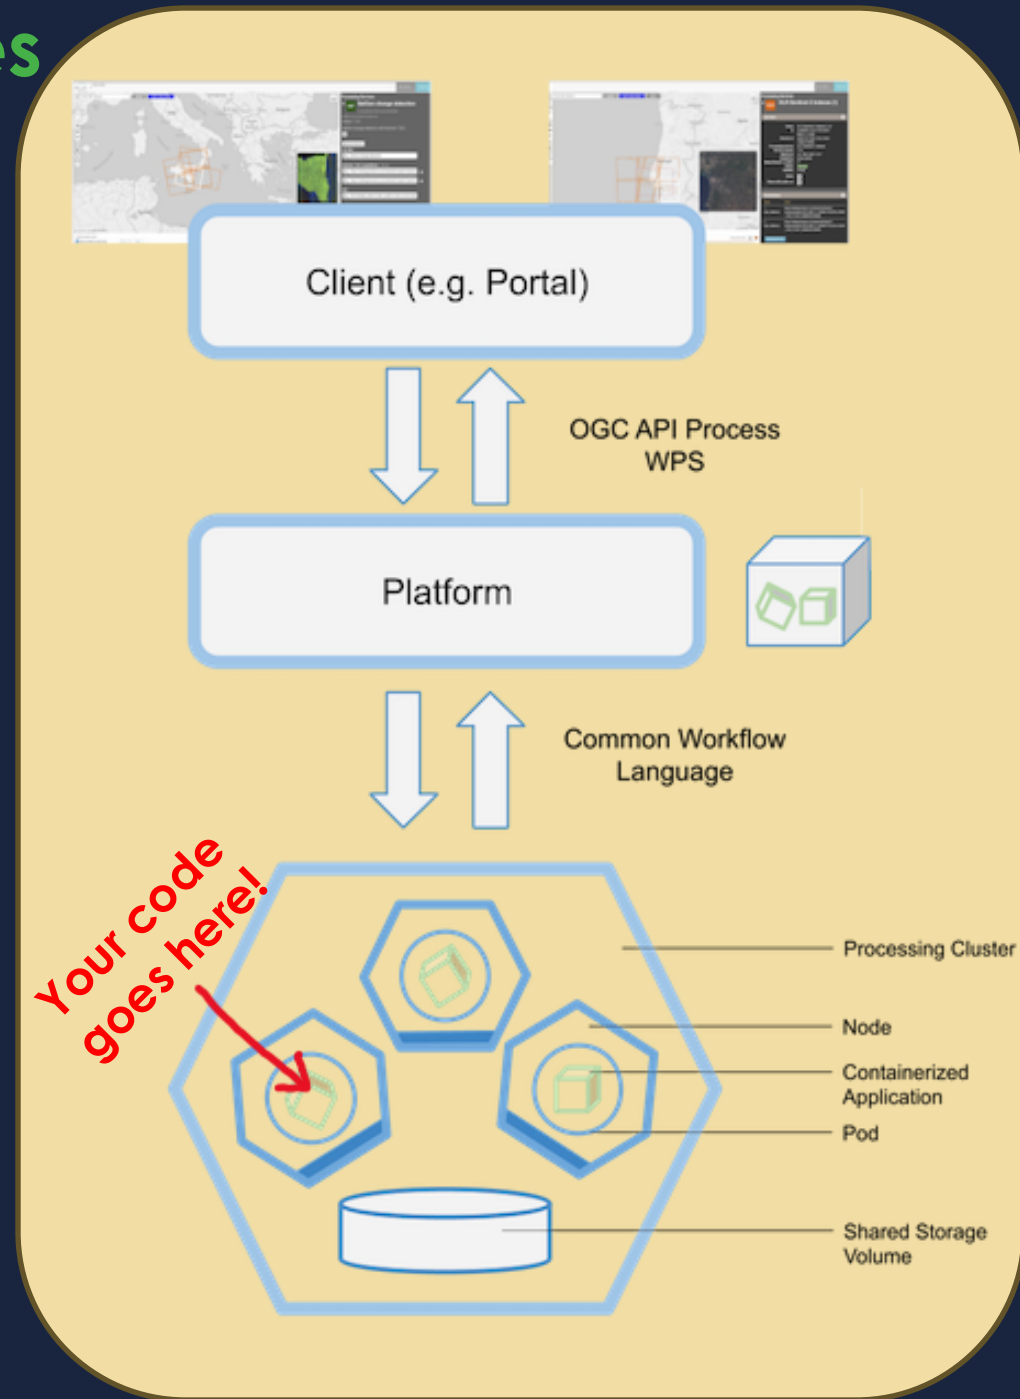
- What's it for?
  - An Application Package provides a **standardized way** to package EO processing software so that it can be run (and combined into workflows) on any compliant platform.
- How does it work? With two components:
  - A **Docker image** which contains the actual code to be run (in any language(s), in any environment).
  - A CWL (**Common Workflow Language**) file which references the Docker image. It details **how to run** the program in the Docker image and what its **inputs and outputs** are.

# Structure of an Application Package



# Running Application Packages

- An Application Package is designed to be run on an **Application Package Platform**.
- For simple local tests, we can use **cwltool**.
- Note: cwltool runs the CWL, but **doesn't** implement stage-in / stage-out like a real platform.

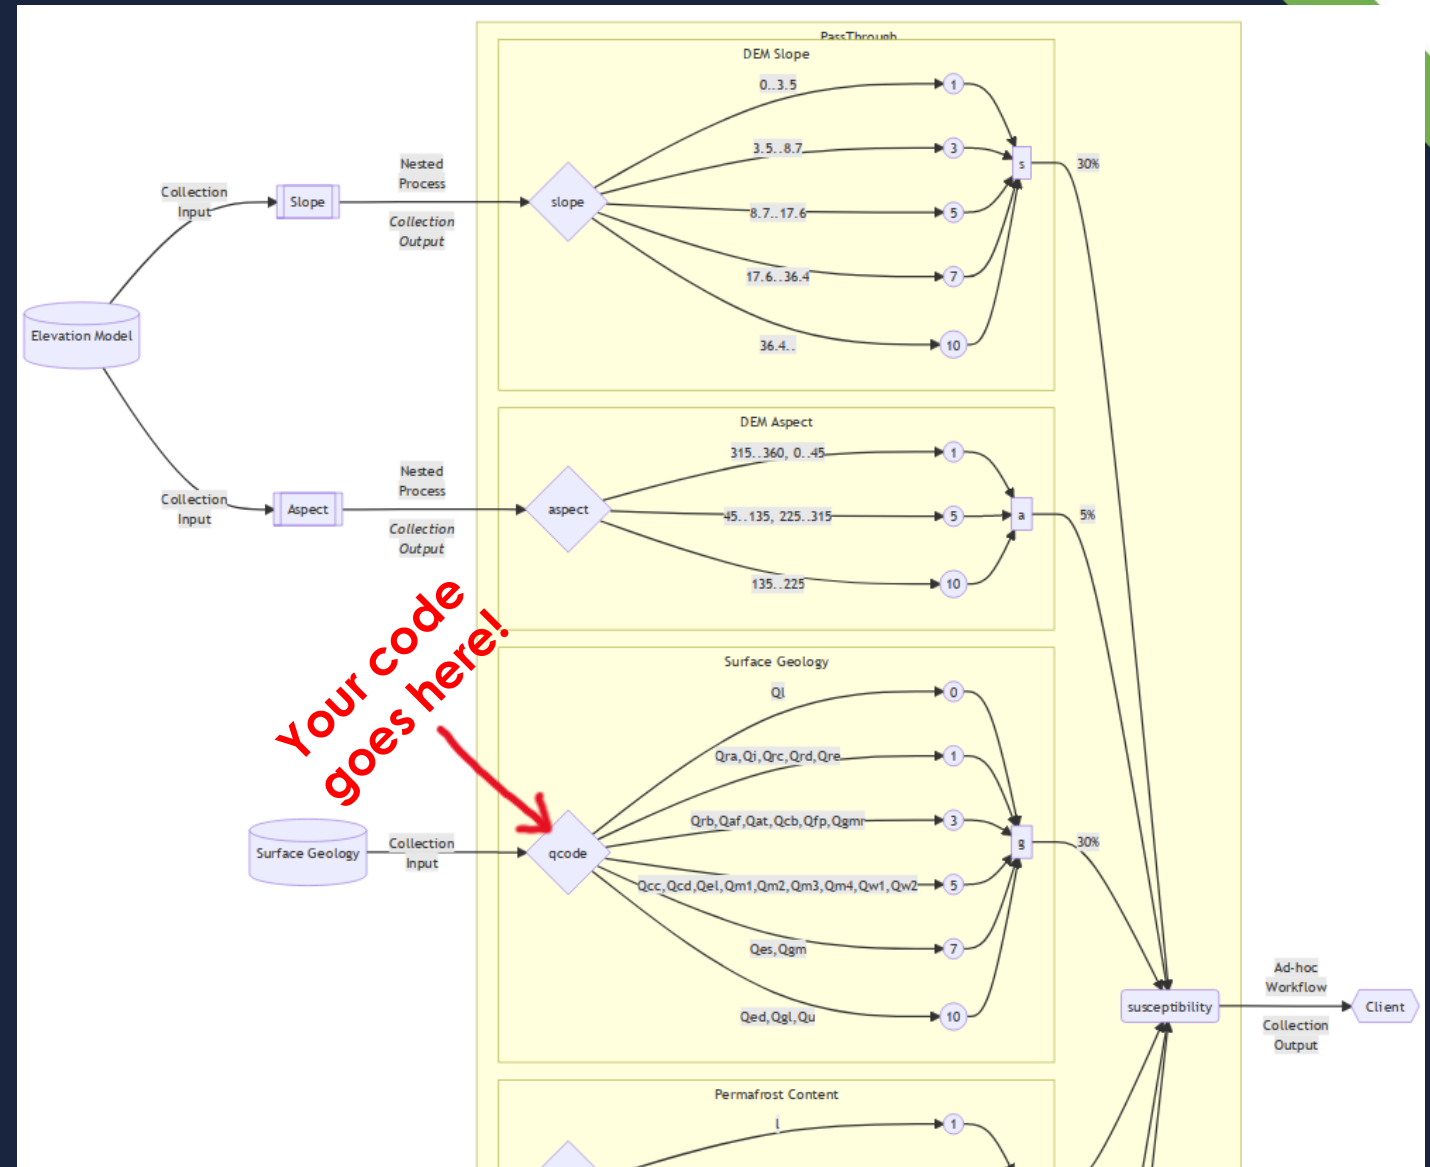




# Application Packages in OGC workflows

An Application Package platform lets you chain and combine Application Packages using OGC API – Processes.

Packaged code becomes a FAIR, versatile, reusable building block.



# What does a CWL file look like?

CWL uses the familiar YAML format, and a simple CWL file has two main parts: a **Workflow** and a **CommandLineTool**.

Here's a simple Workflow with no inputs and one output.

```
- class: Workflow
  id: hello
  label: hello world
  doc: hello world
  requirements: []
  inputs: {}
  outputs:
    - id: stac_catalog
      type: Directory
      outputSource:
        - run_script/results
  steps:
    run_script:
      # References a CommandLineTool
      run: '#myscript'
      in: {}
      out:
        - results
```

# What does a CWL file look like? (continued)

The other main part: the CommandLineTool.

- Runs a specified command in a specified Docker image.
- Inputs/outputs are mapped to the Workflow's inputs/outputs.

```
- class: CommandLineTool
  id: myscript # Referenced by the workflow
  requirements:
    DockerRequirement:
      dockerPull: alpine:3.22.1
    InitialWorkDirRequirement:
      listing:
        - entryname: myscript.sh
          entry: |-
            echo "Hello world!" >>hello.txt
  baseCommand:
    - sh
  arguments:
    - myscript.sh
  inputs: {}
  outputs:
    results:
      type: Directory
      outputBinding:
        glob: .
```



# Hands on 1: edit and run an Application Package



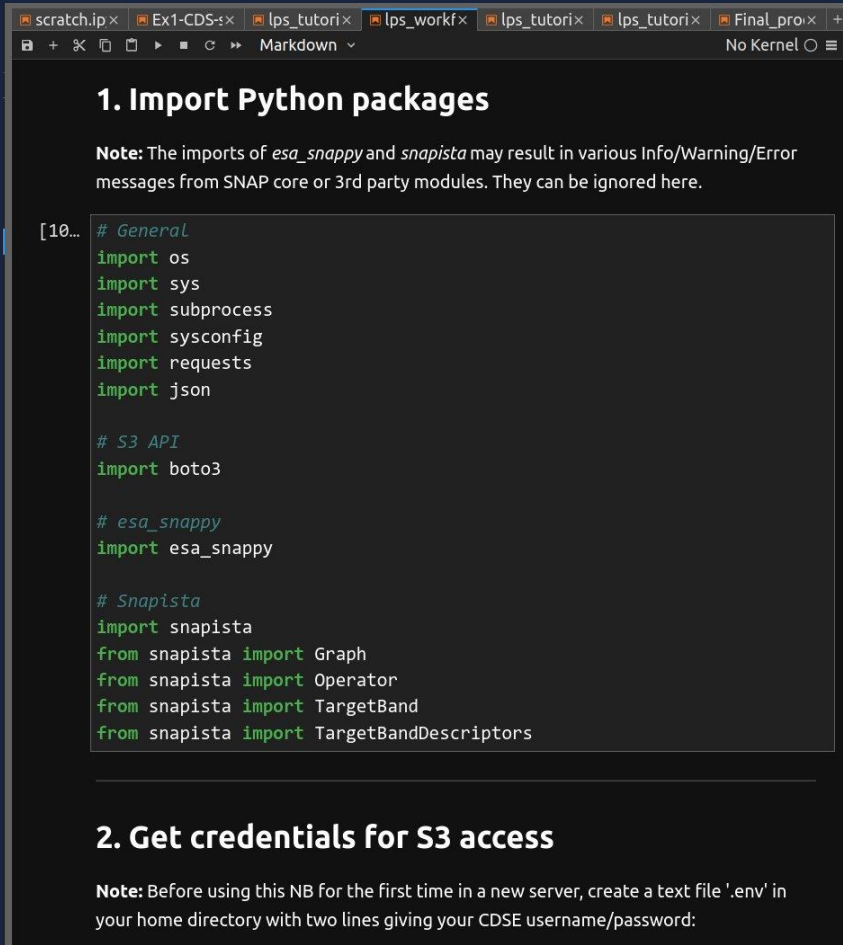
# Part 2: From Notebooks to Application Packages with xcengine

# Comparing Notebooks and Application Packages

Topic	Jupyter Notebook	Application Package
Python environment	Usually handled outside notebook (e.g. conda env file).	Must be set up inside the container.
Parameter handling	Everything's editable, so any variable can be a parameter.	Strictly defined in CWL file, passed to container via CLI.
⚠ Data input/output	Read and write however and wherever you like.	Stage-in and stage-out via platform using STAC catalogues.
Distribution	The notebook file and something defining the environment.	One CWL file referencing one or more container images.
Paradigm	Primarily interactive.	Strictly batch.

# Notebook to Application Package: what's needed?

"Here's my notebook.  
Now what do I do?"



The screenshot shows a Jupyter Notebook with two sections. The first section, '1. Import Python packages', includes a note about potential messages from SNAP core or 3rd party modules and a code cell with imports for os, sys, subprocess, sysconfig, requests, json, boto3, esa\_snappy, and snapista. The second section, '2. Get credentials for S3 access', includes a note about creating a .env file.

```
1. Import Python packages
```

Note: The imports of *esa\_snappy* and *snapista* may result in various Info/Warning/Error messages from SNAP core or 3rd party modules. They can be ignored here.

```
[10... # General
import os
import sys
import subprocess
import sysconfig
import requests
import json

# S3 API
import boto3

# esa_snappy
import esa_snappy

# Snapista
import snapista
from snapista import Graph
from snapista import Operator
from snapista import TargetBand
from snapista import TargetBandDescriptors
```

---

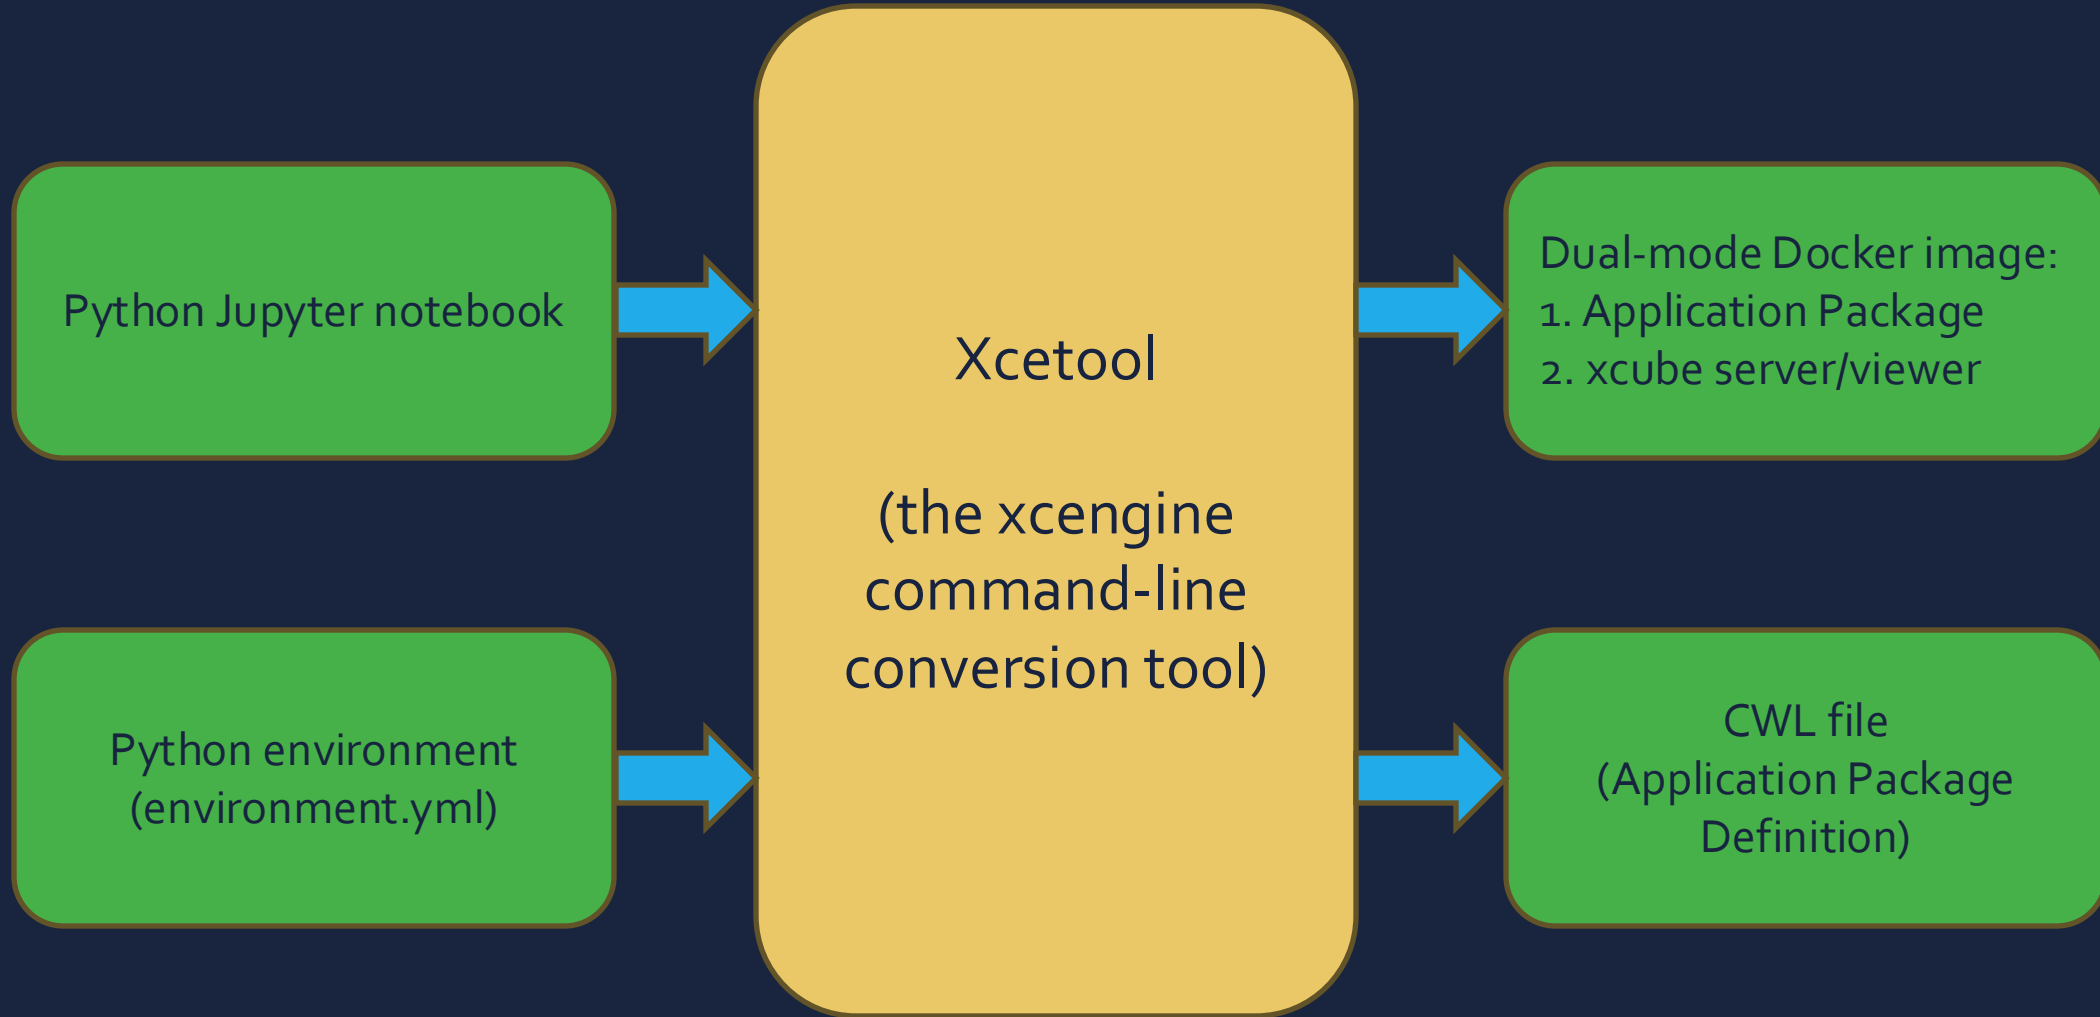
```
2. Get credentials for S3 access
```

Note: Before using this NB for the first time in a new server, create a text file '.env' in your home directory with two lines giving your CDSE username/password:

1. Turn it into a runnable script.
2. Process command-line arguments.
3. Read input data via STAC catalogue.  
(Application Packages **don't** usually fetch their own data!)
4. Write output data via STAC catalogue.
5. Define, build, and publish a container image.
6. Write a CWL file defining inputs/outputs.
7. Test and debug until it works!

*Can we automate some of this?*

# xcengine: turning notebooks into Application Packages



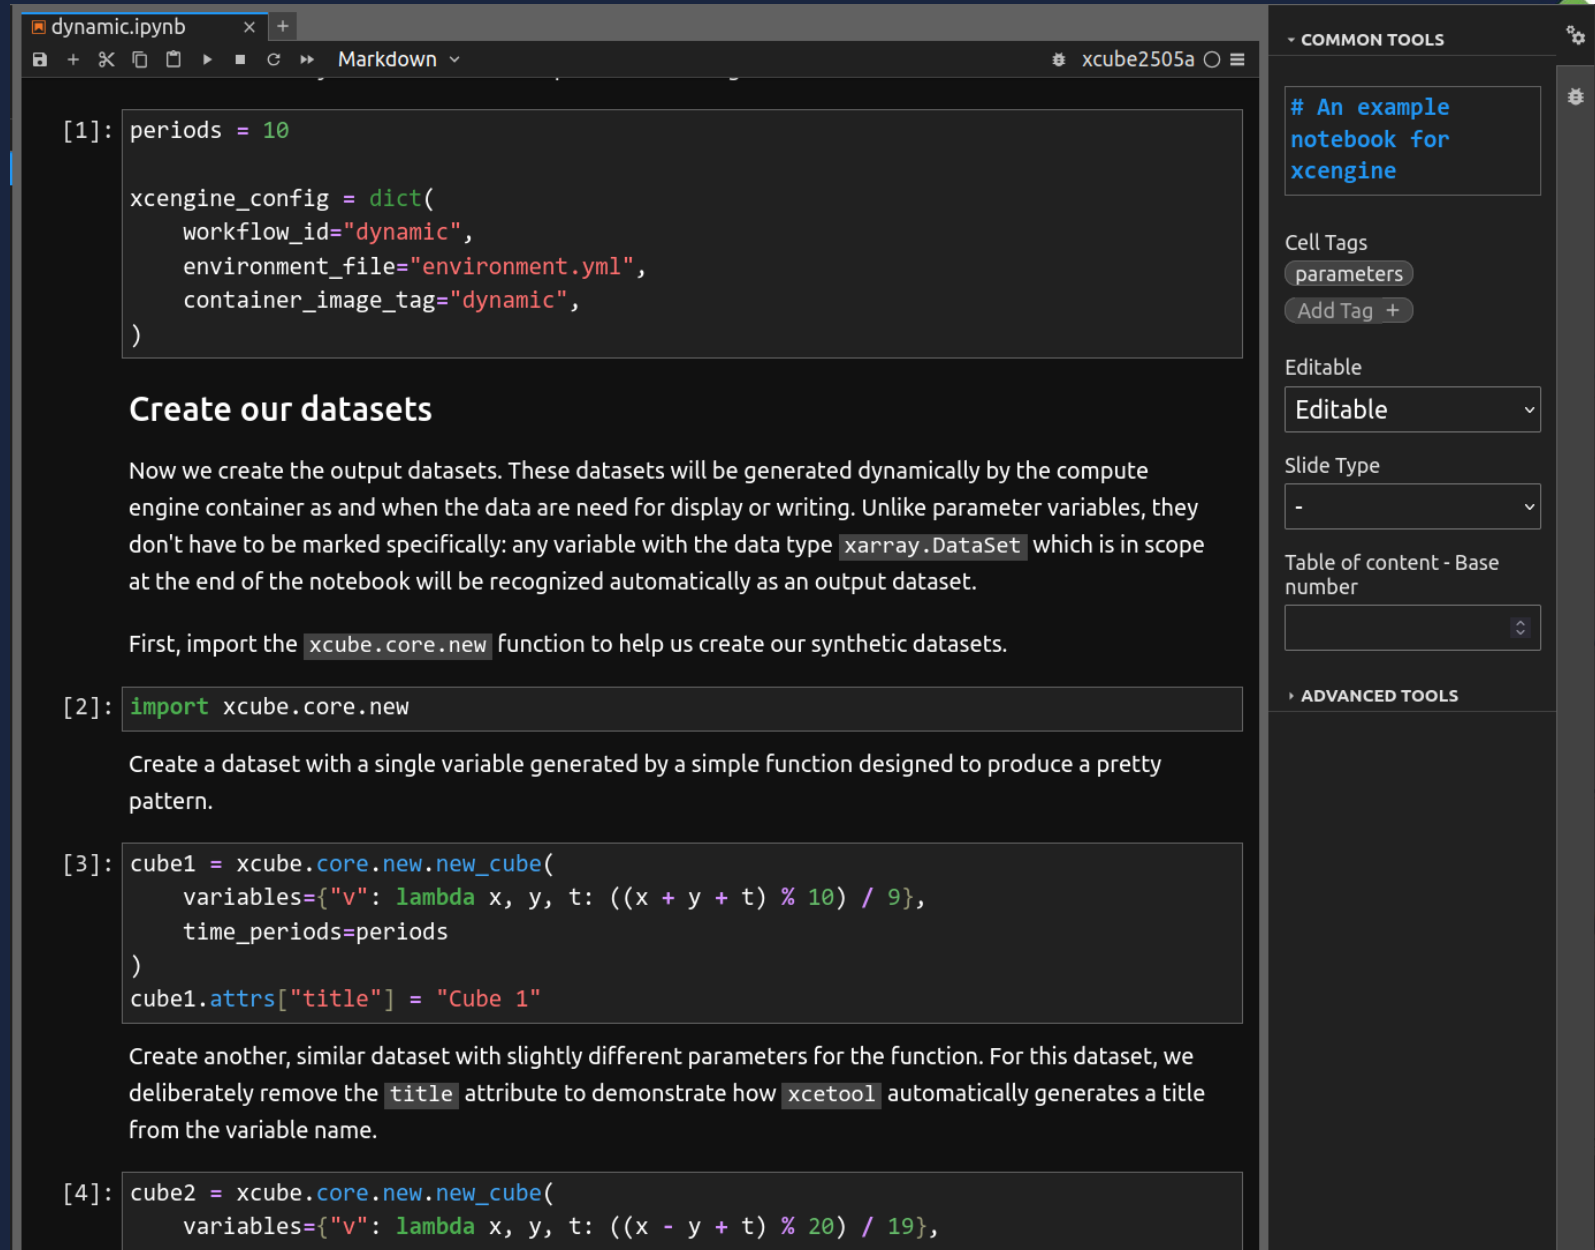




# Hands on 2: use xcengine to create an Application Package

# The input notebook

In the **xcengine-nb** subdirectory of the repository is the notebook **dynamic.ipynb**, which we'll use for an example.



The screenshot displays the Jupyter Notebook interface for a file named `dynamic.ipynb`. The notebook is in `Markdown` view. It contains four code cells and explanatory text.

**Cell [1]:** A code cell defining `xcengine_config` as a dictionary with `workflow_id="dynamic"`, `environment_file="environment.yml"`, and `container_image_tag="dynamic"`.

**Text:** A section titled "Create our datasets" explains that output datasets are generated dynamically by the compute engine container. It notes that variables with the data type `xarray.DataSet` at the end of the notebook will be recognized as output datasets.

**Text:** A paragraph instructs to import the `xcube.core.new` function to create synthetic datasets.

**Cell [2]:** A code cell importing `xcube.core.new`.

**Text:** A paragraph explaining that a dataset with a single variable generated by a simple function designed to produce a pretty pattern.

**Cell [3]:** A code cell creating a dataset `cube1` using `xcube.core.new.new_cube()` with variables `x, y, t` and `time_periods=periods`. The title is set to "Cube 1".

**Text:** A paragraph explaining that another dataset will be created with slightly different parameters, and the `title` attribute is removed to demonstrate how `xcetool` automatically generates a title from the variable name.

**Cell [4]:** A code cell creating a dataset `cube2` using `xcube.core.new.new_cube()` with variables `x, y, t` and `time_periods=periods`. The title is set to "Cube 2".

**Sidebar:** The sidebar on the right contains "COMMON TOOLS" and "ADVANCED TOOLS". Under "COMMON TOOLS", there is a section for "Cell Tags" with a "parameters" tag and an "Add Tag" button. Below that is an "Editable" dropdown menu set to "Editable". Further down is a "Slide Type" dropdown menu set to "-". At the bottom of the sidebar is a "Table of content - Base number" input field.

# The input notebook (continued)

Here's the **parameters cell**, tagged using the Jupyter Lab property inspector. Any variable defined here becomes a parameter.

The screenshot displays the Jupyter Lab interface. On the left, a code cell is visible with the following content:

```
[1]: periods = 10  
  
xcengine_config = dict(  
    workflow_id="dynamic",  
    environment_file="environment.yml",  
    container_image_tag="dynamic",  
)
```



The line `periods = 10` is circled in red, with a red label "variable / property" pointing to it. On the right, the "COMMON TOOLS" sidebar is open. The "Cell Tags" section shows a dropdown menu with "parameters" selected, which is also circled in red with a red label "parameters" tag. Above this, the "property inspector" icon (a gear) is circled in red with a red label "property inspector". Below the "Cell Tags" section, the "Editable" dropdown is set to "Editable", and the "Slide Type" dropdown is set to "-".



# Part 3: xcube in xcengine Docker images

# xcube server and viewer

These are the parts of xcube  
we'll be using today: the built-in  
**API server** and **web viewer**.







# Hands on 3: running the xcengine Docker image in "xcube mode"